Long Ridge Partners Scheduling Agent

This document proposes the design of a scheduling agent for Long Ridge Partners. It is based on our discussions with coordinators and client managers.

This is a first draft; please let us know if there are any mistakes, things you don't like, or corrections. We will work on this until we get it right.

Goals and Non-Goals

Goals

1. **Draft emails for coordinators.** The agent composes scheduling emails including availability requests, confirmations, and Zoom link forwards.
2. **Stay on top of every thread.** The agent keeps track of every in-progress scheduling request and makes sure nothing gets dropped. If a recruiter hasn't responded, a client hasn't confirmed, or a Zoom link hasn't been forwarded, the agent drafts a follow-up for the coordinator.
3. **Update systems automatically.** Upon coordinator approval, the agent creates the shared LRP interview calendar entry, creates the Encore interview activity, and drafts the prep call invite for CM, recruiter, and candidate.

Non-Goals

1. **No autonomous sending.** Every email requires coordinator approval. At this phase, the agent will not even have the capability to send emails.
2. **No visibility beyond coordinators.** Recruiters, CMs, and clients will not interact with or be aware of the agent, although they may get emails from coordinators that the agent drafted.
3. **No candidate-facing communication.** The recruiter owns the candidate relationship. The agent drafts to recruiters, never to candidates.
4. **No judgment calls.** When a situation requires human judgment, the agent asks the coordinator.

Metrics

To ensure a successful engagement, we propose the following metrics to measure business impact:

1. **Mean Time to Interview (MTTI):** Time, in hours, from client interview request to the interview taking place (less is better).
2. **Interview Fulfillment Rate (IFR):** Percentage of client interview requests that result in a completed interview (more is better).

Baseline

Before deploying, we will use data from Encore and Gmail to compute the current metrics over the period Jan-March 2026 so we can measure improvements.

Scheduling flow

The following diagram describes our understanding of the complete scheduling flow, *with* our new AI agent:

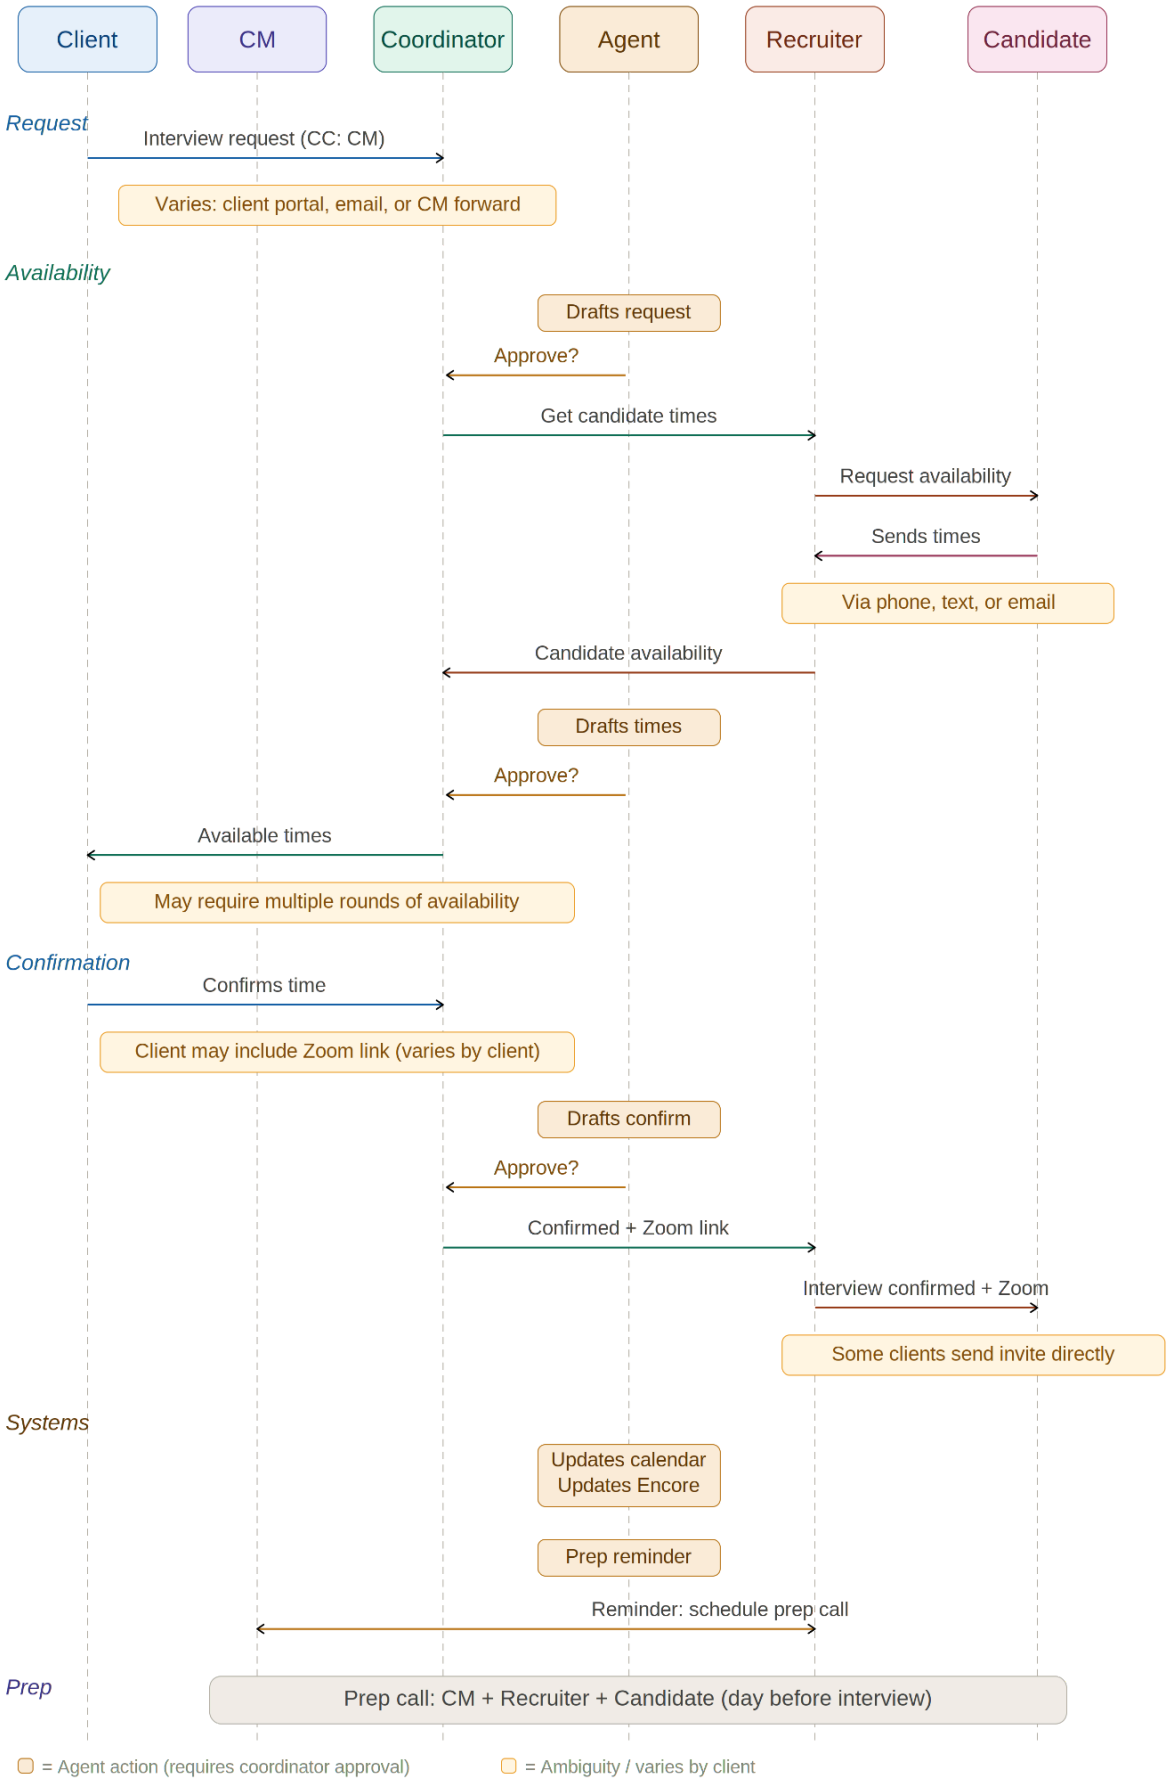
Client
CM
Coordinator
Agent
Recruiter
Candidate
Request
Interview request (CC: CM)
Varies: client portal, email, or CM forward
Availability
Drafts request
Approve?
Get candidate times
Request availability
Sends times
Via phone, text, or email
Candidate availability
Drafts times
Approve?
Available times
May require multiple rounds of availability
Confirmation
Confirms time
Client may include Zoom link (varies by client)
Drafts confirm
Approve?
Confirmed + Zoom link
Interview confirmed + Zoom
Some clients send invite directly
Systems
Updates calendar
Updates Encore
Prep reminder
Reminder: schedule prep call
Prep
Prep call: CM + Recruiter + Candidate (day before interview)
= Agent action (requires coordinator approval)
= Ambiguity / varies by client

User Interface/User Experience

The agent lives inside Gmail as a Google Workspace add-on. It does not require a separate app, new login, or extra browser tab. The add-on has two tabs: drafts and status board.

Drafts tab

- When a coordinator opens a scheduling email, a sidebar panel appears with the agent's recommended action (e.g., a draft to the recruiter requesting availability).
- Three options on every draft: send as-is, edit by hand, or ask for corrections ("make it shorter," "add that we need next week only").
- The agent also creates an "Agent queue" label in Gmail's left sidebar, collecting all active threads with status indicators: drafts ready, threads waiting on responses, and questions needing coordinator judgment.
- The queue is the coordinator's task list. Click in, open a thread, approve in the sidebar, move on.

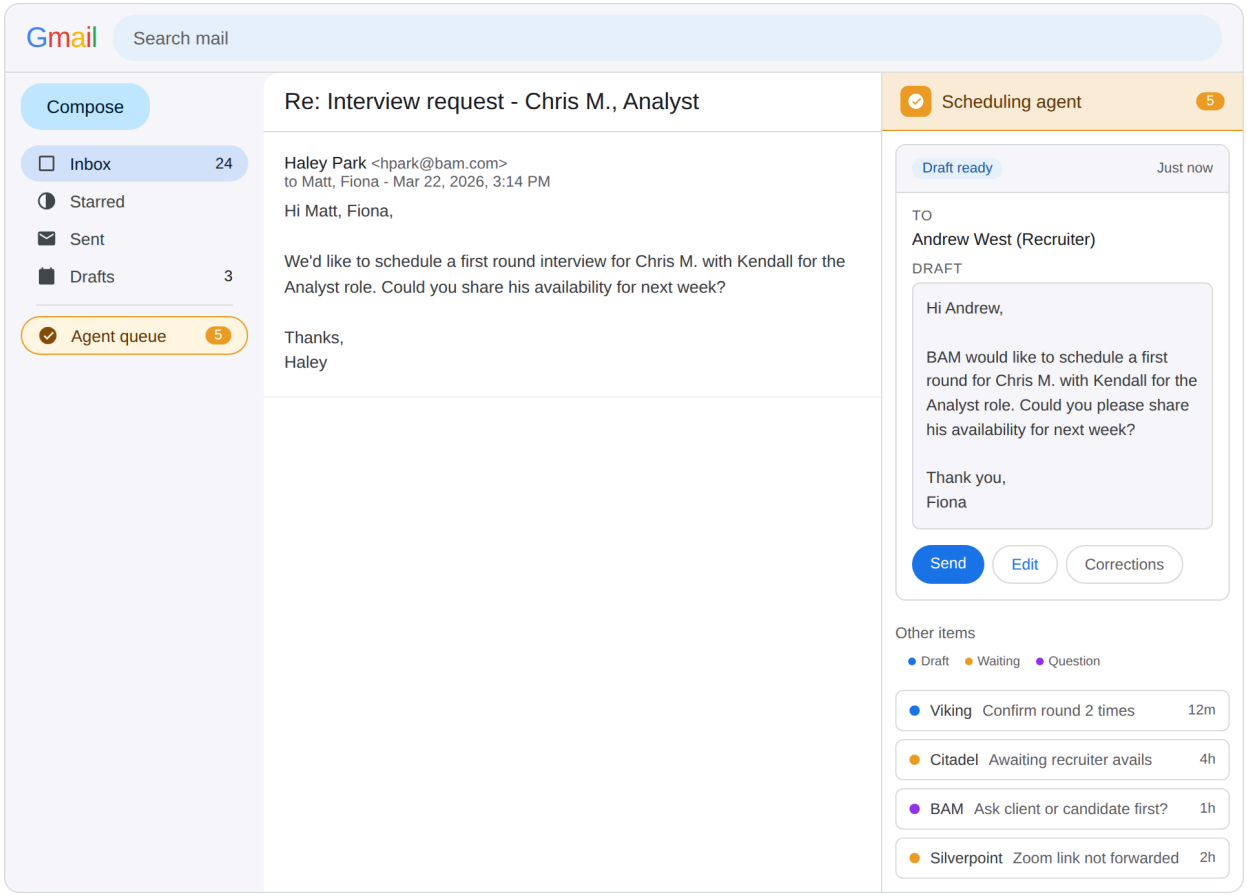

Status board tab

- A second tab in the sidebar shows every active scheduling thread the coordinator manages, grouped by who's blocking it: waiting on recruiter, waiting on client, confirmed but pending systems updates.
- Each thread shows the client, candidate, current step, and how long it's been waiting. Overdue items are flagged in red.
- Summary stats at the top show total active threads, overdue count, and average time to interview.
- This replaces the mental tracking coordinators do today across dozens of threads and hundreds of emails.

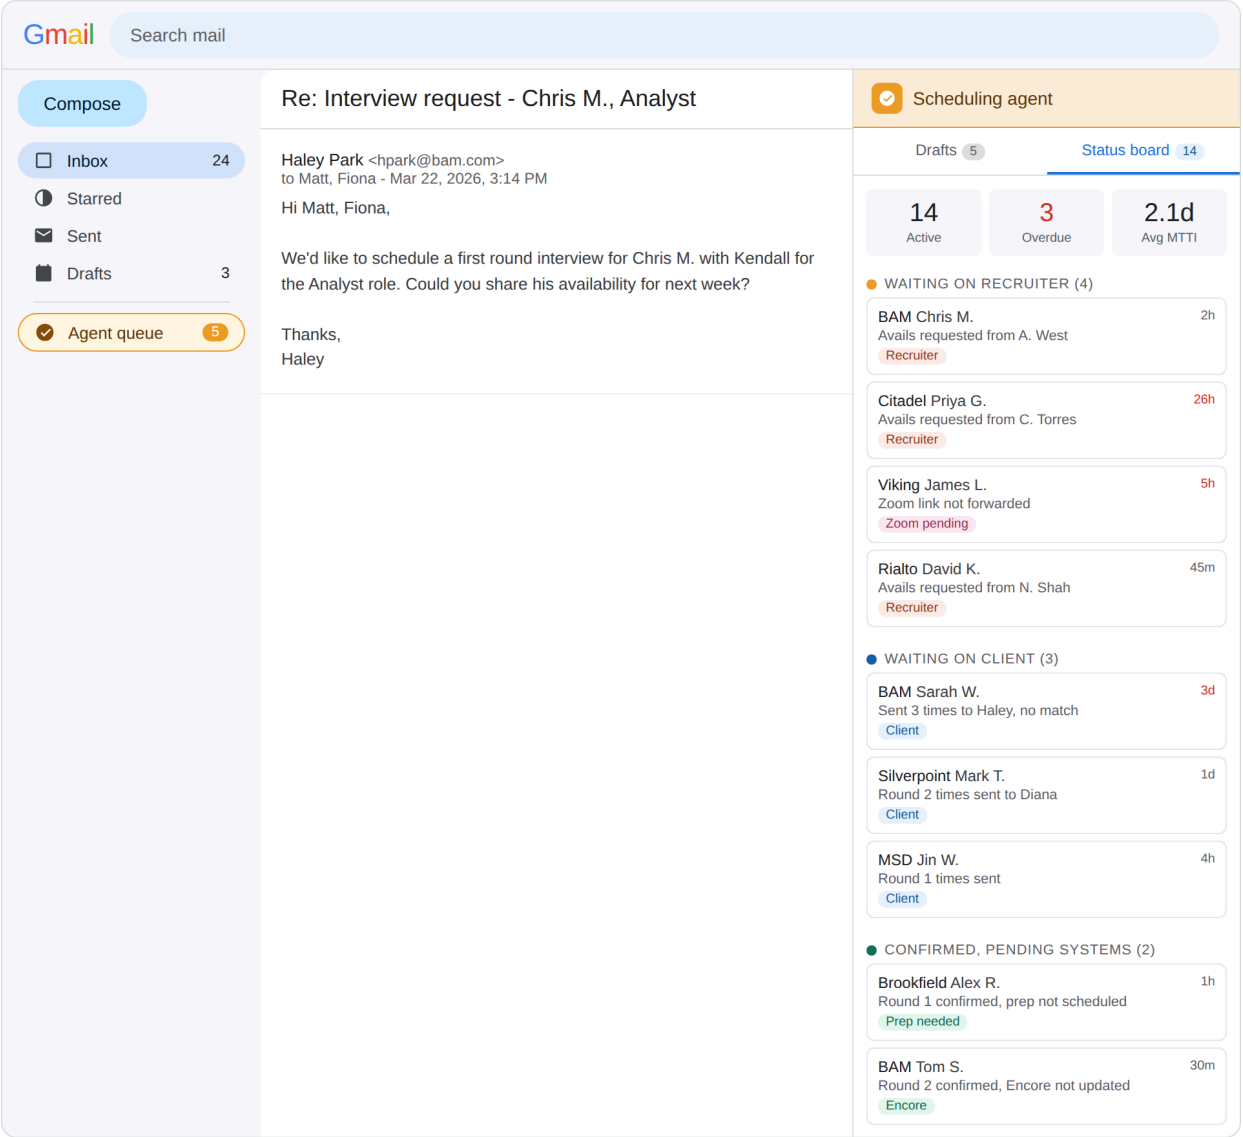

Gmail requirement

The add-on requires the Gmail web interface. It will not work in Outlook or other clients. Since the firm already runs on Google Workspace, this is a UI preference change for one coordinator rather than a platform migration. We'd like to confirm this is acceptable.

Learning

It learns client-specific workflows on the fly. When it’s not sure what to do, it asks the coordinator, and it saves the response to its internal client profile. Over time, it learned the bespoke communications patterns of each LRP client.

System architecture

The agent runs as a backend service hosted in the cloud. It connects to three systems:

- **Gmail** (via Google Workspace APIs): Reads incoming scheduling emails, creates draft responses, sends approved emails on behalf of coordinators, and manages the Agent queue label. Uses domain-wide delegation, meaning the agent accesses coordinator accounts through a secure service account authorized by LRP's Google Workspace admin. It does not have access to any accounts beyond the coordinators.
- **Google Calendar** (via Google Workspace APIs): Creates and updates entries on the shared LRP interview calendar and schedules prep call invites. Same authorization mechanism as Gmail.
- **Encore** (via Cluein Data Connector): Reads candidate and recruiter records to identify the correct parties in a scheduling thread. Writes interview activities when interviews are confirmed. Access will be established through Cluein's support team, starting with a read/write connection scoped to interview-related data only.

The Gmail sidebar add-on is the only user-facing piece. It calls the backend service when a coordinator opens a thread, receives the agent's recommendation, and sends the coordinator's response (approve, edit, correction) back to the backend.

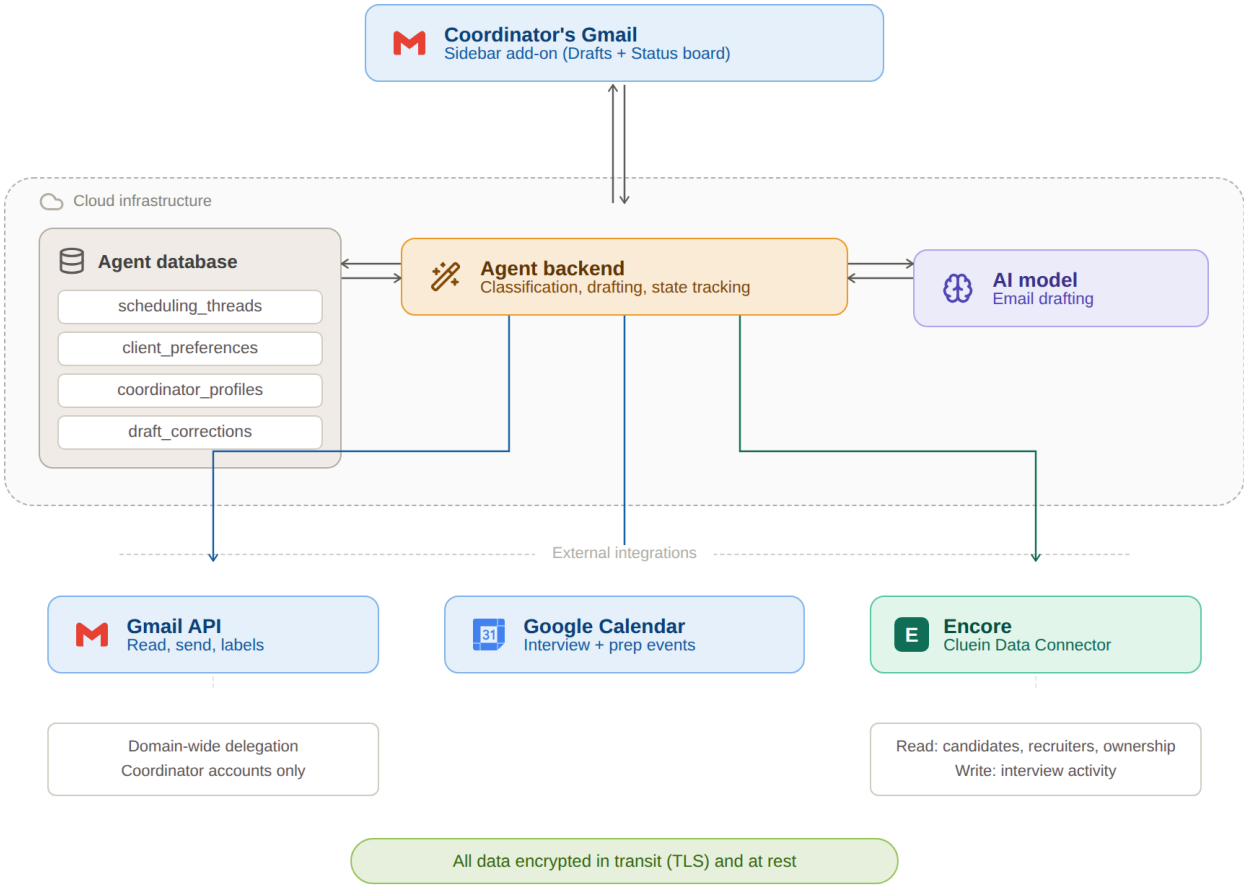

Data model

The agent maintains a small internal database tracking:

- **Scheduling threads**: One record per active interview request. Tracks current state (awaiting availability, awaiting confirmation, confirmed, etc.), which parties are involved, timestamps at each stage, and the next expected action.
- **Client preferences**: Learned per-client settings like Zoom link behavior, who sends calendar invites, and whether to ask the client or candidate for availability first. Updated when coordinators answer sidebar questions or correct the agent.
- **Coordinator profiles**: Which coordinator owns which clients, and any draft corrections they've made that inform future drafts.

No candidate resumes, compensation data, or other sensitive recruiting information is stored in the agent's database. Candidate identification uses email addresses as unique keys, consistent with how Encore and Quill operate.

How a request flows through the system

1. A scheduling email arrives in the coordinator's Gmail inbox.

2. The backend detects it, classifies it (new request, availability response, confirmation, reschedule, cancellation), and identifies the relevant client, candidate, and recruiter from the thread and Encore records.
3. The backend determines the next action and drafts a response.
4. When the coordinator opens the email, the sidebar displays the draft and action buttons.
5. The coordinator approves (or edits/corrects). The backend sends the email from the coordinator's account, updates the thread state, and triggers any system updates (calendar entry, Encore activity, prep call invite).

If the coordinator doesn't open the email, the item sits in the Agent queue with escalating age indicators. The agent does not act without approval.

Where learning happens

- **Client preferences** are stored when coordinators answer the agent's first-time questions and when they correct a wrong assumption. This is explicit, not inferred.
- **Draft quality** improves when coordinators use the "Corrections" button. The agent stores what was asked for and applies it to similar future drafts for that client.
- No model fine-tuning or retraining occurs. Learning is stored as structured preferences and rules, not as changes to the underlying AI model.

Data security

- All data is encrypted in transit (TLS) and at rest.
- The agent's Gmail access is scoped exclusively to coordinator accounts. It cannot read recruiter, CM, or executive email.
- Encore access is scoped to interview-related records only.
- The agent's internal database contains no sensitive candidate information beyond names and email addresses needed for scheduling.
- No data is shared with third parties. The AI model provider (Anthropic) does not retain or train on any data processed through the agent.
- LRP's Google Workspace admin retains full control over the agent's access and can revoke it at any time through the admin console.

Key risks

- **Wrong candidate.** The agent identifies candidates by email address, not name. When multiple Encore records match or no email is present, it asks the coordinator rather than guessing.
- **Wrong recruiter.** If the thread's recruiter doesn't match the Encore owner field, the agent surfaces the discrepancy instead of silently picking one.
- **Overlapping schedules.** The agent tracks candidate availability across all coordinator accounts and flags time conflicts before confirmation.
- **Misdirected email.** Every email requires explicit coordinator approval. The agent has no autonomous send capability.

- **System sync failures.** If Encore or Calendar is unreachable, the agent queues the update and surfaces the failure in the status board.
- **Agent downtime.** The agent is an overlay, not a replacement. If it goes down, coordinators schedule manually as they do today. Nothing breaks.